

Title: Scalable Cross-Attention Mechanisms

Abstract: We present a novel transformer architecture that computes self-attention with linear complexity.

Section 2: Methodology

The proposed model processes multi-head embeddings efficiently across GPU clusters.