

IEEE Paper: Scalable Vision Transformers

Abstract: We tackle the high computational cost of self-attention in vision transformers. Standard attention scales

Section 3: Linear Cross-Attention Architecture

We propose Linear-X Attention that replaces standard dot-product attention with kernel feature maps, reducing c

Section 4: Experiments on ImageNet-1k

Our model achieves 84.8% top-1 accuracy on ImageNet-1k benchmark, outperforming Swin-Transformer while u